

Dani Roytburg

droytbur@andrew.cmu.edu
djroytburg.github.io
github.com/djroytburg · linkedin.com/in/danijroytburg

October 4, 2026

Summary Master’s student in the Machine Learning Department at Carnegie Mellon University focused on reliable and interpretable AI, multi-agent training paradigms, and activation-based model analysis.

Education

M.S.	Carnegie Mellon University, School of Computer Science Machine Learning	2027
B.S.	Emory University, College of Arts and Science Computer Science; Quantitative and Social Sciences	2025

Positions

ML Alignment and Theory Scholars Berkeley, CA	Research Fellow) May 2026-September 2026
---	---

Advised by Shi Feng. Developed persona-training methods which emulate midtraining for better model organisms and constitutions. Engineered evaluations across 100+ training runs on synthetic-document pipelines. Evaluated how tastefully frontier models can test ideas in safety research.

Carnegie Mellon University Pittsburgh, PA	Research Assistant, DI Group September 2025-present
---	--

Develop multi-agent reinforcement learning environments to improve legibility of reasoning traces. Design user studies and adversarial setups to measure influence of legibility on human and multi-agent interactions.

Martian Research San Francisco, CA	Research Fellow May 2025-February 2026
--	---

Engineered activation profiles of language model evaluators to suppress self-preferential bias, applied to LLM routers (ICML 2026, MechInterp @ NeurIPS 2025). Winner of Mechanistic Router Interpretability Hackathon (40 submissions).

Emory University Atlanta, GA	Research Assistant, Digital Humanities Lab November 2021-May 2025
--	--

Designed and trained large-scale graph-text models; evaluated adversarial robustness of LLMs; built historical correspondence networks and dashboards.

ZS Associates
Evanston, IL

Internship - Business Technology Solutions
June 2024-August 2024

Automated ETL pipelines and implemented synthetic-data imputation for compromised client data.

Cloverpop
Chicago, IL

Intern — Machine Learning
January 2024-May 2024

Deployed decision-management systems and structured-prediction models on meeting transcripts.

McCormick and Company
Hunt Valley, MD

Intern — Global Marketing and Sales Data Science
June 2023-August 2023

Built Amazon sales data visualizations and automated analytics with GCP BigQuery and BERT-based embedding search.

Northeastern University
Boston, MA

Research Assistant, Mediacloud Lab
June 2022-October 2022

Extended NER to long-tail ethnics cuisines using Bon Appetit data.

Publications

Conference Papers

Dani Roytburg, Matthew Bozoukov, Matthew Nguyen, Jou Barzdukas, Mackenzie Puig-Hall, and Narmeen Oozeer. Are llm evaluators really narcissists? sanity checking self-preference evaluations. In *Forty-Third International Conference on Machine Learning*, 2026. URL: <https://arxiv.org/abs/2601.22548>

Dani Roytburg, Shreya Sridhar, and Daphne Ippolito. Measuring weak-to-strong legibility of reasoning models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2026. Also presented at: ICML 2026 Workshop on AI for Good (AI4GOOD). URL: <https://openreview.net/forum?id=IR4nD1PK12>

Dani Roytburg and Daphne Ippolito. Disentangling models from personas in heterogeneous llm simulations. In *Second Workshop on Social Simulation with LLMs: Fidelity in Applications (COLM 2026, Spotlight)*, 2026. URL: <https://openreview.net/forum?id=ZBC0tbXyhg>

Dani Roytburg and Beck Miller. Mind the gap: Pathways towards unifying ai safety and ethics research. In *Proceedings of the International Association for Safe and Ethical AI*, 2026. URL: <https://arxiv.org/abs/2512.10058>

Dani Roytburg, Matthew Bozoukov, Hongyu Fu, Matthew Nguyen, Jou Barzdukas, and Narmeen Fatimah Oozeer. Breaking the mirror: Activation-based mitigation of self-preference in llm evaluators. In *Mechanistic Interpretability Workshop at NeurIPS 2025*, 2025. Also accepted at: NeurIPS 2025 Workshop on Evaluating the Evolving LLM Lifecycle; NeurIPS 2025 Workshop on Reliable ML from Unreliable Data. URL: <https://arxiv.org/abs/2509.03647>

Dani Roytburg, Deborah Olorunisola, Sandeep Soni, and Lauren Klein. Words and action: Modeling linguistic leadership in # blacklivesmatter communities. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 19, pages 1690–1703, 2025. URL: <https://ojs.aaai.org/index.php/ICWSM/article/view/35895>

Thesis

Recognition

Awards

Best Poster, LTI Student Research Symposium, 2026

Emory University Dean's List, 2022, 2024-2025

Winner, Martian Research Mechanistic Interpretability Hackathon, 2025

Skills

Python Java R JavaScript Typescript SQL PyTorch JAX HuggingFace Transformers scikit-learn
spaCy networkx D3.js React.js GCP Docker MySQL