

DANI ROYTBURG

droytbur@andrew.cmu.edu | [djroytburg.github.io](https://github.com/droytburg) | [in/danijroytburg](https://in.danijroytburg)

Education

Carnegie Mellon University, School of Computer Science

Pittsburgh, PA

M.S., Machine Learning

(on leave) expected May 2027

Relevant Coursework: Advanced Machine Learning, Probability and Statistics, Probabilistic Graphical Models, Machine Learning in Practice

Service: Executive Board Member, CMU AI Safety Initiative.

Emory University, College of Arts and Science

Atlanta, GA

B.S., *summa cum laude*: Computer Science, Quantitative and Social Sciences.

August 2021-May 2025

Cumulative GPA: 3.85/4.0 | **Dean's List** Spring 2024

Relevant Coursework: Deep Learning, Machine Learning, Privacy-Enhanced Machine Learning, Systems Programming, Computer Architectures, Linear Algebra, Adv. Calculus, Probability and Statistics (I, II), Discrete Math, Algorithmic Analysis.

Work Experience

ML Alignment and Theory Scholars

Berkeley, CA

Research Fellow, advised by [Shi Feng](#)

May 2026 – present

- Develop persona-training methods which emulate midtraining for better model organisms and constitutions.
- Engineer evaluations across 100+ training runs on synthetic-document pipelines.
- Released an [independent evaluation](#) of frontier models' ability to propose tasteful safety research.

DI Group, Carnegie Mellon University

Pittsburgh, PA

Research Assistant, advised by [Daphne Ippolito](#)

September 2025 – present

- Develop multi-agent reinforcement learning environments to improve legibility of reasoning traces.
- Design user studies and adversarial setups to measure influence of legibility on human and multi-agent interactions.

Martian Research

San Francisco, CA

Research Fellow, advised by [Jacob Haimes](#), [Phillip Quirke](#), Narmeen Oozeer

May 2025 – February 2026

- Engineer activation profiles of language model evaluators to suppress self-preferential bias, applied to LLM routers (ICML 2026, MechInterp @ NeurIPS 2025).
- Winner of [Mechanistic Router Interpretability](#) Hackathon (40 submissions).

Digital Humanities Lab, Emory University

Atlanta, GA

Research Assistant, advised by [Lauren Klein](#) and [Sandeep Soni](#)

November 2021 – May 2025

- Designed and trained terabyte-scale graph-text models of linguistic network propagation (ICWSM '25).
- Evaluated adversarial robustness of large language models' commonsense reasoning to distributional shift.
- Built historical correspondence networks and analytic dashboards in (D3, Next.js), supported by PBS.

ZS Associates

Evanston, IL

Intern – Business Technology Solutions Associate

June 2024 – August 2024

- Automated ETL pipelines from diverse data sources for consumption forecasts of specialty neurological treatments.

Select Publications

[Measuring Weak-to-Strong Legibility](#) of Reasoning Models (2026)

Dani Roytburg, Shreya Sridhar, Daphne Ippolito. *Neural Information Processing Systems (NeurIPS 2026)*

[Are LLM Evaluators Really Narcissists?](#) Sanity Checking Self-Preference Evaluations (2026)

Dani Roytburg, Narmeen Oozeer et. al, *International Conference on Machine Learning (ICML 2026)*

[Breaking the Mirror](#): Activation-Based Mitigation of Self-Preference in LLM Evaluators (2025)

Dani Roytburg, Narmeen Oozeer et. al, *Mechanistic Interpretability, Eval-LLM Workshops @ NeurIPS 2025*

[Mind the Gap](#): Pathways Towards Unifying AI Safety and Ethics Research (2026)

Dani Roytburg, Beck Miller, *International Association for Safe & Ethical AI (IASEAI) 2026*

[Words and Action](#): Modeling Linguistic Leadership in #BlackLivesMatter (2025)

Dani Roytburg, Debbie Olorunisola, Sandeep Soni, Lauren Klein, *AAAI ICWSM 2025*

Technical Skills

Languages: Python, Java, R, JavaScript, Typescript, SQL || French (C1 proficiency), Russian (C1), Chinese (B1)

Packages: PyTorch, JAX, HuggingFace Transformers, scikit-learn, multiprocessing, spaCy, networkx, D3.js

Databases and Deployment: Firebase, Vercel, Amazon EMR, Google Cloud Platform (GCP), Docker, MySQL

Dev Tools: Git, Linux, Power BI, Tableau, Figma